




# How to set up an OpenTelemetry Collector (OTEL) in a Kubernetes cluster

Table of contents

Issue

Environment

Cause

Resolution

## Issue

This setup consists of three main components:

- OpenTelemetry Collector Configuration:** Defines the behavior of the Collector, including the receivers, processors, and exporters.
- OpenTelemetry Collector Deployment:** The actual deployment of the OpenTelemetry Collector within Kubernetes.
- Fluent Bit Configuration:** Configures Fluent Bit to forward logs to the OpenTelemetry Collector.

## Environment

Ezmeral Unified Analytics Version: 1.3

## Cause

This article provides a step-by-step guide on how to set up an OpenTelemetry Collector in a Kubernetes cluster. The goal is to receive logs from Fluent Bit and forward them to another OpenTelemetry host.

## Resolution

### 1. OpenTelemetry Collector Configuration

The first step involves creating a ConfigMap to store the OpenTelemetry Collector configuration. This configuration includes:

- OTLP receiver:** To listen on port 4318 for HTTP requests.
- Batch processor:** For performance optimization.
- OTLP exporter:** To forward data to another OpenTelemetry host.

Below is the configuration for the ConfigMap:

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: otel-collector-config
  namespace: monitoring
data:
  otel-collector-config.yaml: |
    receivers:
      otlp:
        protocols:
          http:
            endpoint: "0.0.0.0:4318"

    processors:
      batch:

    exporters:
      otlp:
        endpoint: "another-otel-host:4317"
        tls:
          insecure: true # Set to false and configure certs for production use

    service:
      pipelines:
        logs:
          receivers: [otlp]
          processors: [batch]
          exporters: [otlp]
```

Key Points:

- **Receiver (otlp)**: Accepts logs via HTTP on port 4318.
- **Processor (batch)**: Optimizes log forwarding by batching data.
- **Exporter (otlp)**: Forwards logs to the specified endpoint (another OpenTelemetry host). In production, you should enable proper TLS configuration by setting `insecure` to `false`.

## 2. OpenTelemetry Collector Deployment

Next, we deploy the OpenTelemetry Collector using a Kubernetes `Deployment` resource. This deployment will:

- Use the official OpenTelemetry Collector Docker image.
- Mount the configuration from the previously created ConfigMap.
- Expose port 4318 for receiving logs.

Below is the configuration for the deployment and service:

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: otel-collector
  namespace: monitoring
  labels:
    app: otel-collector
spec:
  replicas: 1
  selector:
    matchLabels:
      app: otel-collector
  template:
    metadata:
      labels:
        app: otel-collector
    spec:
      containers:
        - name: otel-collector
          image: otel/opentelemetry-collector:0.72.0
          args:
            - "--config=/conf/otel-collector-config.yaml"
          ports:
            - containerPort: 4318
          volumeMounts:
            - name: otel-collector-config
              mountPath: /conf
      volumes:
        - name: otel-collector-config
          configMap:
            name: otel-collector-config
---
apiVersion: v1
kind: Service
metadata:
  name: otel-collector
  namespace: monitoring
spec:
  selector:
    app: otel-collector
  ports:
    - protocol: TCP
      port: 4318
      targetPort: 4318
```

### Key Points:

- **Kubernetes Deployment**: Deploys the OpenTelemetry Collector with the configuration mounted from the ConfigMap.
- **Kubernetes Service**: Exposes the OpenTelemetry Collector within the cluster to allow Fluent Bit to send logs.

## 3. Fluent Bit Configuration

Now, we update the Fluent Bit configuration to send logs to the OpenTelemetry Collector. This setup includes:

- The OpenTelemetry output plugin.
- Filters to specifically send logs from pods with the label `hpe-ezua/app=ezpresto`.

Here's the Fluent Bit configuration:

```
[SERVICE]
    flush 1
    Log_Level info
    Daemon off
    HTTP_Server On
    HTTP_Listen 0.0.0.0
    HTTP_Port 2020
    Parsers_File parsers.conf
    health_check On

[INPUT]
    Name tail
    Tag kube.*
    Path /var/log/containers/*.log
    Parser docker
    Mem_Buf_Limit 50MB
    Buffer_Max_Size 1MB
    Skip_Long_Lines Off
    Read_From_Head true
    Refresh_Interval 10

[FILTER]
    Name kubernetes
    Match kube.*
    Merge_Log On
    Keep_Log On
    Labels On

# Output to OpenTelemetry Collector for ezpresto pods
[OUTPUT]
    Name opentelemetry
    Match kube.*
    Host otel-collector.monitoring.svc.cluster.local
    Port 4318
    Metrics_uri /v1/metrics
    Logs_uri /v1/logs
    Tls On
    Tls.verify Off

<Rule>
    Key_exists $kubernetes['labels']['hpe-ezua/app']
    Key_value_equals $kubernetes['labels']['hpe-ezua/app'] ezpresto
</Rule>
```

### Key Points:

- **Input Section:** Fluent Bit collects logs from container files and system logs.
- **Filters:** Refines logs and includes metadata like Kubernetes labels.
- **Output Section:** Sends logs matching specific criteria (pods with `hpe-ezua/app=ezpresto`) to the OpenTelemetry Collector.

To apply the new configuration, update the Fluent Bit ConfigMap and restart the Fluent Bit DaemonSet:  
`kubectl -n monitoring rollout restart daemonset fluentbit`

## Configuration Adjustments

Before deploying, ensure the following adjustments are made:

- **Replace `another-otel-host`** in the configuration with your actual OpenTelemetry host.
- **Adjust the namespace** if necessary.
- **Update OpenTelemetry image version** as required.
- **Enable proper TLS** for production environments by configuring certificates and setting `insecure` to `false`.

## Deployment Steps

1. **Create and apply the ConfigMap** with the OpenTelemetry Collector configuration:  
`kubectl apply -f otel-collector-config.yaml`
2. **Create and apply the Deployment and Service** for the OpenTelemetry Collector:  
`kubectl apply -f otel-collector-deployment.yaml`

3. **Update the Fluent Bit configuration** to send logs to the OpenTelemetry Collector. Then, restart Fluent Bit for the new configuration to take effect:  
`kubectl -n monitoring rollout restart daemonset fluentbit`

## Conclusion

This setup creates a log processing pipeline where:

- Fluent Bit collects logs from Kubernetes pods.
- Logs are sent to the OpenTelemetry Collector deployed in the cluster.
- The OpenTelemetry Collector forwards these logs to another OpenTelemetry host.

This architecture allows for flexible and scalable log collection and forwarding within a Kubernetes environment, streamlining observability and log management.

